

# Security Properties for Stack Safety

ANONYMOUS AUTHOR(S)

The call stack is a perennial target for low-level attacks, leading to consequences that range from leakage or corruption of private stack data to control-flow hijacking. A profusion of software and hardware protections have been proposed for detecting or preventing such attacks, including stack canaries [2], bounds checking [5, 9, 10], split stacks [7], shadow stacks [3, 12], capabilities [1, 6, 13–15], and hardware tagging [11]. The protections offered by such mechanisms are commonly described in terms of concrete examples of attacks that they can prevent. At best, stack safety is defined by reference to an idealized machine that is arguably stack safe by construction [14]. But these mechanisms can be intricate, and it would be useful to have a precise, generic, and formal specification for stack safety, both to compare the security claims of different enforcement techniques and to rigorously validate such claims.

We propose such a characterization using the tools of language-based security. The informal claim that “stack safety protects a caller from its callee” amounts to saying that it guarantees the *integrity* and *confidentiality* of the caller’s local state until it regains control. We formalize these integrity and confidentiality requirements as trace (hyper-)properties, with two different variants: *stepwise* variants, in which a caller’s data must *never* be read or modified during a call, and *observational* ones, in which callees may read from and write to their caller’s stack frame, as long as these “risky” behaviors do not affect the system’s observable behavior. The observational properties are more extensional and thus allow “lazier” protection mechanisms that permit risky reads or writes as long as they do not affect visible outcomes.

Confidentiality is especially interesting. It is based on a traditional notion of noninterference, but whereas ordinary noninterference is an end-to-end hyper-property on whole program runs, stack confidentiality is a “nested” form of noninterference applied to subtraces of the whole program trace delimited by call instructions and corresponding returns, requiring that the callee’s behavior is invariant under hypothetical changes to [BCP: scrambling of?] the caller’s stack frame.

This formulation not only captures the intuition that the callee cannot directly access the caller’s state; it also gives a novel way of looking at control-flow attacks. We are not interested in [BCP: I find “We are not interested in” confusing here—the previous sentence promises a novel way of looking, and this opening feels irrelevant to that.] attacks that merely aim to execute arbitrary code – indeed, our model assumes that attackers (i.e., callees) may already execute arbitrary code. By extension, we do not place any particular restrictions on control flow. Perhaps suprisingly, we can still meaningfully distinguish the caller from its callee, and talk about protecting the caller’s data. After a call, control is considered to belong to the callee until the trace reaches a valid “return target,” typically a state with the program counter at the next instruction and the stack pointer restored. If the caller’s data are accessed before then, it is a violation

of integrity or confidentiality. [BCP: Actually, the whole paragraph doesn’t make very much sense to me.]

These properties can optionally be extended with a notion of *well-bracketed control flow* as in Skorstengaard et al. [14]—a global requirement that callees must always return to their immediate caller, if they return at all. This property is largely [BCP: “largely” feels like a weasel word...] orthogonal to our data-protection properties.

To demonstrate the utility of our properties, we use them to evaluate an existing mechanism, the *stack-safety micro-policies* of Roessler and DeHon [11], re-implemented in the Coq proof assistant on top of a RISC-V ISA model [BCP: instead of “spec”. OK?]. We use QuickChick [4, 8], a property-based testing tool for Coq, to generate random programs and check that Roessler and Dehon’s micro-policies correctly abort programs that would violate stack safety. Our testing regime is also able to detect incorrect variants of the enforcement mechanisms—both variants that we accidentally created during our re-implementation of the micro-policy and ones that we intentionally crafted to be broken in order to increase our confidence in testing and the enforcement mechanism itself. For these broken mechanisms, our tester generates counterexamples that violate our properties but escape the enforcement. This increases our confidence in enforcement mechanisms for which the tester is unable to find counterexamples.

Our testing supports the stack-safety claim of Roessler and Dehon’s *Depth Isolation* micro-policy, in which memory cells within each stack frame are tagged with the depth of the function activation that owns the frame and access to those locations is then permitted only when that activation is currently executing. On the other hand, we find that their *Lazy Tagging and Clearing* policy violates the temporal aspect of confidentiality in corner cases where data can leak across repeated calls to the same callee, and also violates integrity if the leak happens to use the caller’s frame. We propose a variant of *Lazy Tagging and Clearing* that should [BCP: “that we believe will?”] enforce observational integrity and confidentiality, albeit at some performance cost. (An obvious next step is to expand testing to these properties and validate the variant, but efficiently testing observational properties is more challenging because a property violation can be arbitrarily far removed from the initial “risky” read or write.)

We have demonstrated [BCP: let’s use the same verb tense consistently—“we do X” [my preference] or “we have done X” or “we did X”—all the way through] our model’s flexibility by extending it to different settings. First, we can support passing arguments on the stack, allowing some of the state of the caller to be shared with the callee (and potentially with a sub-callee, and so on.) We can also describe a simple coroutine model, with one statically bounded stack per coroutine, each protected from the others.

In ongoing work, we are investigating whether stack protection schemes based on the Cheri capability system [1, 15] can enforce our properties. In particular, we are adapting the “uninitialized capability

model” of Georges et al. [6] to work with our testing framework in order to validate it against our properties. Out of several Cheri-based stack protection schemes, this one’s treatment of uninitialized frames seems most compatible with our model of confidentiality, without the need for expensive stack clearing between calls. We expect that Cheri-based techniques can support our properties if we assume extra cooperation from callers, namely that they do not leak their capabilities to the attacker.

## REFERENCES

- [1] David Chisnall, Colin Rothwell, Robert N.M. Watson, Jonathan Woodruff, Munraj Vadera, Simon W. Moore, Michael Roe, Brooks Davis, and Peter G. Neumann. 2015. Beyond the PDP-11: Architectural Support for a Memory-Safe C Abstract Machine. In *Proceedings of the Twentieth International Conference on Architectural Support for Programming Languages and Operating Systems* (Istanbul, Turkey) (ASPLOS ’15). Association for Computing Machinery, New York, NY, USA, 117–130. <https://doi.org/10.1145/2694344.2694367>
- [2] Crispin Cowan, Calton Pu, Dave Maier, Heather Hintony, Jonathan Walpole, Peat Bakke, Steve Beattie, Aaron Grier, Perry Wagle, and Qian Zhang. 1998. StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks. In *Proceedings of the 7th Conference on USENIX Security Symposium - Volume 7* (San Antonio, Texas) (SSYM’98). USENIX Association, USA, 5.
- [3] Thurston H.Y. Dang, Petros Maniatis, and David Wagner. 2015. The Performance Cost of Shadow Stacks and Stack Canaries. In *Proceedings of the 10th ACM Symposium on Information, Computer and Communications Security* (Singapore, Republic of Singapore) (ASIA CCS ’15). Association for Computing Machinery, New York, NY, USA, 555–566. <https://doi.org/10.1145/2714576.2714635>
- [4] Maxime Dénès, Catalin Hritcu, Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C. Pierce. 2014. QuickChick: Property-Based Testing for Coq (abstract). In VSL. <http://www.easychair.org/smart-program/VSL2014/index.html>
- [5] Joe Devietti, Colin Blundell, Milo M. K. Martin, and Steve Zdancewic. 2008. Hard-Bound: Architectural Support for Spatial Safety of the C Programming Language. In *13th International Conference on Architectural Support for Programming Languages and Operating Systems*. 103–114. [http://acg.cis.upenn.edu/papers/asplos08\\_hardbound.pdf](http://acg.cis.upenn.edu/papers/asplos08_hardbound.pdf)
- [6] Aïna Linn Georges, Armaël Guéneau, Thomas Van Strydonck, Amin Timany, Alix Trieu, Sander Huyghebaert, Dominique Devriese, and Lars Birkedal. 2021. Efficient and provable local capability revocation using uninitialized capabilities. *Proc. ACM Program. Lang.* 5, POPL (2021), 1–30. <https://doi.org/10.1145/3434287>
- [7] Volodymyr Kuznetsov, László Szekeres, Mathias Payer, George Candea, R. Sekar, and Dawn Song. 2014. Code-Pointer Integrity. In *Proceedings of the 11th USENIX Conference on Operating Systems Design and Implementation* (Broomfield, CO) (OSDI’14). USENIX Association, USA, 147–163.
- [8] Leonidas Lampropoulos and Benjamin C. Pierce. 2018. *QuickChick: Property-Based Testing in Coq*. Electronic textbook. Version 1.0. <http://www.cis.upenn.edu/~bcpierce/sf>.
- [9] Santosh Nagarakatte, Jianzhou Zhao, Milo M. K. Martin, and Steve Zdancewic. 2009. SoftBound: highly compatible and complete spatial memory safety for C. In *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 245–258. [http://repository.upenn.edu/cgi/viewcontent.cgi?article=1941&context=cis\\_reports](http://repository.upenn.edu/cgi/viewcontent.cgi?article=1941&context=cis_reports)
- [10] Santosh Nagarakatte, Jianzhou Zhao, Milo M. K. Martin, and Steve Zdancewic. 2010. CETS: compiler enforced temporal safety for C. In *9th International Symposium on Memory Management*. ACM, 31–40. [http://acg.cis.upenn.edu/papers/ismm10\\_cets.pdf](http://acg.cis.upenn.edu/papers/ismm10_cets.pdf)
- [11] Nick Roessler and André DeHon. 2018. Protecting the Stack with Metadata Policies and Tagged Hardware. In *2018 IEEE Symposium on Security and Privacy, SP 2018, Proceedings, 21-23 May 2018, San Francisco, California, USA*. IEEE Computer Society, 478–495. <https://doi.org/10.1109/SP.2018.00066>
- [12] Vedvyas Shanhogue, Deepak Gupta, and Ravi Sahita. 2019. Security Analysis of Processor Instruction Set Architecture for Enforcing Control-Flow Integrity. In *Proceedings of the 8th International Workshop on Hardware and Architectural Support for Security and Privacy* (Phoenix, AZ, USA) (HASP ’19). Association for Computing Machinery, New York, NY, USA, Article 8, 11 pages. <https://doi.org/10.1145/3337167.3337175>
- [13] Lau Skorstengaard, Dominique Devriese, and Lars Birkedal. 2019. Reasoning about a Machine with Local Capabilities: Provably Safe Stack and Return Pointer Management. *ACM Trans. Program. Lang. Syst.* 42, 1, Article 5 (Dec. 2019), 53 pages. <https://doi.org/10.1145/3363519>
- [14] Lau Skorstengaard, Dominique Devriese, and Lars Birkedal. 2019. StkTokens: Enforcing Well-Bracketed Control Flow and Stack Encapsulation Using Linear Capabilities. *Proc. ACM Program. Lang.* 3, POPL, Article 19 (Jan. 2019), 28 pages. <https://doi.org/10.1145/3290332>
- [15] Jonathan Woodruff, Robert N.M. Watson, David Chisnall, Simon W. Moore, Jonathan Anderson, Brooks Davis, Ben Laurie, Peter G. Neumann, Robert Norton, and Michael Roe. 2014. The Cheri Capability Model: Revisiting RISC in an Age of Risk. In *Proceeding of the 41st Annual International Symposium on Computer Architecture* (Minneapolis, Minnesota, USA) (ISCA ’14). IEEE Press, 457–468.